

Project 03 – Applications

Arthur J. Redfern

arthur.redfern@utdallas.edu

Nov 11, 2020

1 Logistics

- Assigned Nov 11, 2020 and due Dec 02, 2020
- This is an individual project, no help from others is allowed
- You cannot choose a project that is similar to a project you've done previously (in another class, for fun, for work, ...); if the project I select for you falls into this category you need to alert me immediately and a different project will be selected
- The use of any and all online resources is allowed
- All text should be written by you – do not simply copy text from others

2 Goals

- This class covered a lot, from fundamental math to network design, training and implementation to a variety of applications; you have a strong basis to work from
- In this project you will take a deep dive into a topic that we didn't cover in detail in this class, demonstrating your ability to learn and summarize new topics

3 Project

Rank the topics below based on your interest and immediately send me an email with your 1 (most interested) – N (least interested) list. I'll reply back with the project that I've selected for you. While this is an individual project, I'm limiting the number of people that can select the same topic.

If you have an appropriate topic that you're very interested in that's not on this list, there's a small possibility I'll allow it. Send me an email with the topic, key references and your justification along with the above 1 – N list for backup and again, I'll let you know the project I've selected for you.

After receiving your project selection confirmation from me create a ≤ 20 slide tutorial that includes:

- Title slide
- Motivation slide
- Brief summary of previous work / background slide(s)
- Topic summary covering key ideas slides
- References (not included in the ≤ 20 slide limit)

Limit the number of equations to just the key ones and maximize the number of figures and plots used to illustrate the key ideas. Text should come from you, not be copied from others. If you copy a figure from somewhere else, provide a reference.

I'm not looking for a disjoint regurgitation of 5 separate papers. I'm looking for someone who understands an area and who is able to create a coherent summary of the area built from understanding and shaping multiple key papers / ideas into a clear presentation / story arc. An interesting story takes the reader on a journey that is both exciting and logically flowing. Think about the story your slides tell.

The resulting slides should enable a person who is generally skilled in the area of xNNs (e.g., 1 of your classmates or me) to learn the basics of a new topic. I like slides that a person can read offline and understand and also work in the context of an oral presentation (even though we're not going to do oral presentations). Make your slides with this in mind as I'm the 1 doing the grading.

The following are the 11 potential project topics to be ranked by you 1 – N in your email to me. Emails that do not contain a ranking of all N will be ignored. Key references to start with are provided for all cases, but feel free to expand your search as appropriate. You don't need to include material from all the references if it doesn't fit into the story you tell.

3.1 Topic 1: Design / Conditional Computation And Dynamic Networks

WeightNet: revisiting the design space of weight networks

<https://arxiv.org/abs/2007.11823>

DyNet: dynamic convolution for accelerating convolutional neural networks

<https://arxiv.org/abs/2004.10694>

Dynamic region-aware convolution

<https://arxiv.org/abs/2003.12243>

Conditional convolutions for instance segmentation

<https://arxiv.org/abs/2003.05664>

CondConv: conditionally parameterized convolutions for efficient inference

<https://arxiv.org/abs/1904.04971>

Squeeze-and-excitation networks

<https://arxiv.org/abs/1709.01507>

3.2 Topic 2: Design / Graph Neural Networks

Deep graph library: towards efficient and scalable deep learning on graphs

<https://arxiv.org/abs/1909.01315>

<https://github.com/dmlc/dgl>

Fast graph representation learning with PyTorch geometric

<https://arxiv.org/abs/1903.02428>

https://github.com/rusty1s/pytorch_geometric

A comprehensive survey on graph neural networks

<https://arxiv.org/abs/1901.00596>

Graph neural networks: a review of methods and applications

<https://arxiv.org/abs/1812.08434>

How powerful are graph neural networks?

<https://arxiv.org/abs/1810.00826>

Stanford CS224W: machine learning with graphs

See the slides from lectures 8, 9, 10, 18 and 19

<http://web.stanford.edu/class/cs224w/>

<https://colab.research.google.com/drive/1DIQm9rOx2mT1bZETEeVUThxcrP1RKqAn>

Graph deep learning

<https://towardsdatascience.com/graph-deep-learning/home>

The graph neural network model

<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.1015.7227&rep=rep1&type=pdf>

An introduction to graph neural networks: models and applications

https://www.youtube.com/watch?v=zCEYiCxrL_0

3.3 Topic 3: Training / Modern ImageNet Training Optimizations

Fixing the train-test resolution discrepancy: FixEfficientNet

<https://arxiv.org/abs/2003.08237>

Fixing the train-test resolution discrepancy

<https://arxiv.org/abs/1906.06423>

Big transfer (BiT): general visual representation learning

<https://arxiv.org/abs/1912.11370>

Adversarial examples improve image recognition

<https://arxiv.org/abs/1911.09665>

Self-training with noisy student improves ImageNet classification

<https://arxiv.org/abs/1911.04252>

Billion-scale semi-supervised learning for image classification

<https://arxiv.org/abs/1905.00546>

Bag of tricks for image classification with convolutional neural networks

<https://arxiv.org/abs/1812.01187>

AutoAugment: learning augmentation policies from data

<https://arxiv.org/abs/1805.09501>

Mixup: beyond empirical risk minimization

<https://arxiv.org/abs/1710.09412>

3.4 Topic 4: Implementation / Binary Neural Networks

Training binary neural networks with real-to-binary convolutions

<https://arxiv.org/abs/2003.11535>

ReActNet: towards precise binary neural network with generalized activation functions

<https://arxiv.org/abs/2003.03488>

MeliusNet: can binary neural networks achieve MobileNet-level accuracy?

<https://arxiv.org/abs/2001.05936>

TentacleNet: a pseudo-ensemble template for accurate binary convolutional neural networks

<https://arxiv.org/abs/1912.10103>

XNOR-Net++: improved binary neural networks

<https://arxiv.org/abs/1909.13863>

Structured binary neural networks for image recognition

<https://arxiv.org/abs/1909.09934>

MoBiNet: a mobile binary network for image classification

<https://arxiv.org/abs/1907.12629>

Structured binary neural networks for accurate image classification and semantic segmentation

<https://arxiv.org/abs/1811.10413>

Bi-Real Net: enhancing the performance of 1-bit CNNs with improved representational capability and advanced training algorithm

<https://arxiv.org/abs/1808.00278>

Towards accurate binary convolutional neural network

<https://arxiv.org/abs/1711.11294>

XNOR-Net: ImageNet classification using binary convolutional neural networks

<https://arxiv.org/abs/1603.05279>

3.5 Topic 5: Implementation / Graph Compilers

The deep learning compiler: a comprehensive survey

<https://arxiv.org/abs/2002.03794>

TVM: an automated end-to-end optimizing compiler for deep learning

<https://arxiv.org/abs/1802.04799>

Apache TVM (incubating) an end to end deep learning compiler stack for CPUs, GPUs and accelerators

<https://tvm.apache.org>

TVM documentation

<https://tvm.apache.org/docs/>

2019 TVM and deep learning compilation conference: morning keynote & session 1

<https://www.youtube.com/watch?v=npqO0hVXZwU>

3.6 Topic 6: Vision / Anchor Box Free Detection And Segmentation

SOLOv2: dynamic, faster and stronger

<https://arxiv.org/abs/2003.10152>

SOLO: segmenting objects by locations

<https://arxiv.org/abs/1912.04488>

<https://github.com/aim-uofa/AdelaiDet/>

Bridging the gap between anchor-based and anchor-free detection via adaptive training sample selection

<https://arxiv.org/abs/1912.02424>

<https://github.com/sfzhang15/ATSS>

CenterMask : real-time anchor-free instance segmentation

<https://arxiv.org/abs/1911.06667>

<https://github.com/youngwanLEE/CenterMask>

NAS-FCOS: fast neural architecture search for object detection

<https://arxiv.org/abs/1906.04423>

Objects as points

<https://arxiv.org/abs/1904.07850>

<https://github.com/xingyizhou/CenterNet>

FCOS: fully convolutional one-stage object detection

<https://arxiv.org/abs/1904.01355>

<https://github.com/tianzhi0549/FCOS>

3.7 Topic 7: Vision / Point Cloud Processing

Deep learning for 3D point clouds: a survey

<https://arxiv.org/abs/1912.12033>

A convolutional decoder for point clouds using adaptive instance normalization

<https://arxiv.org/abs/1906.11478>

Learning object bounding boxes for 3D instance segmentation on point clouds

<https://arxiv.org/abs/1906.01140>

FlowNet3D: learning scene flow in 3d point clouds

<https://arxiv.org/abs/1806.01411>

VoxelNet: end-to-end learning for point cloud based 3D object detection

<https://arxiv.org/abs/1711.06396>

PointNet++: deep hierarchical feature learning on point sets in a metric space

<https://arxiv.org/abs/1706.02413>

PointNet: deep learning on point sets for 3D classification and segmentation

<https://arxiv.org/abs/1612.00593>

3.8 Topic 8: Vision / Transformers

An image is worth 16x16 words: transformers for image recognition at scale

<https://arxiv.org/abs/2010.11929>

https://github.com/google-research/vision_transformer

<https://github.com/lucidrains/vit-pytorch>

https://www.youtube.com/watch?v=TrdevFK_am4

End-to-end object detection with transformers

<https://arxiv.org/abs/2005.12872>

Exploring self-attention for image recognition

<https://arxiv.org/abs/2004.13621>

Adaptive attention span in computer vision

<https://arxiv.org/abs/2004.08708>

On the relationship between self-attention and convolutional layers

<https://arxiv.org/abs/1911.03584>

Stand-alone self-attention in vision models

<https://arxiv.org/abs/1906.05909>

Image transformer

<https://arxiv.org/abs/1802.05751>

3.9 Topic 9: Language / Large Scale Language Models

Hugging face transformers

<https://github.com/huggingface/transformers>

<https://huggingface.co/transformers/>

Linformer: self-attention with linear complexity

<https://arxiv.org/abs/2006.04768>

Reformer: the efficient transformer

<https://arxiv.org/abs/2001.04451>

BART: denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension

<https://arxiv.org/pdf/1910.13461.pdf>

Exploring the limits of transfer learning with a unified text-to-text transformer

<https://arxiv.org/abs/1910.10683>

DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter

<https://arxiv.org/abs/1910.01108>

ALBERT: a lite BERT for self-supervised learning of language representations

<https://arxiv.org/abs/1909.11942>

RoBERTa: a robustly optimized BERT pretraining approach

<https://arxiv.org/abs/1907.11692>

XLNet: generalized autoregressive pretraining for language understanding

<https://arxiv.org/abs/1906.08237>

Transformer-XL: attentive language models beyond a fixed-length context

<https://arxiv.org/abs/1901.02860>

BERT: pre-training of deep bidirectional transformers for language understanding

<https://arxiv.org/abs/1810.04805>

3.10 Topic 10: Speech / Text To Speech

ESPnet-TTS: unified, reproducible, and integratable open source end-to-end text-to-speech toolkit

<https://arxiv.org/abs/1910.10909>

<https://github.com/espnet/espnet>

The LJ speech dataset

<https://keithito.com/LJ-Speech-Dataset/>

LibriTTS: a corpus derived from LibriSpeech for text-to-speech

<https://arxiv.org/abs/1904.02882>

<http://www.openslr.org/60/>

FastSpeech 2: fast and high-quality end-to-end text to speech

<https://arxiv.org/abs/2006.04558>

FastSpeech: fast, robust and controllable text to speech

<https://arxiv.org/abs/1905.09263>

Neural speech synthesis with transformer network

<https://arxiv.org/abs/1809.08895>

Transfer learning from speaker verification to multispeaker text-to-speech synthesis

<https://arxiv.org/abs/1806.04558>

Style tokens: unsupervised style modeling, control and transfer in end-to-end speech synthesis

<https://arxiv.org/abs/1803.09017>

Natural TTS synthesis by conditioning wavenet on mel spectrogram predictions

<https://arxiv.org/abs/1712.05884>

Tacotron: towards end-to-end speech synthesis

<https://arxiv.org/abs/1703.10135>

WaveNet: a generative model for raw audio

<https://arxiv.org/abs/1609.03499>

3.11 Topic 11: Games / StarCraft II And Or Dota2

3.11.1 StarCraft II

AlphaStar: Grandmaster level in StarCraft II using multi-agent reinforcement learning

<https://deepmind.com/blog/article/AlphaStar-Grandmaster-level-in-StarCraft-II-using-multi-agent-reinforcement-learning>
https://www.nature.com/articles/s41586-019-1724-z.epdf?author_access_token=IzH3nqPYtWJXfDA10WOCNNRgN0jAjWel9jnR3ZoTv0PSZcPzJFGNAZhOlK4deBCKzKm70KfinloafEF1bCCXL6IIHHgKaDkaTkBcTEv7aT-wqDoG1VeO9-wO3GEoAMF9bAOt7mJ0RWQnRVMbyfgH9A%3D%3D

AlphaStar: Mastering the Real-Time Strategy Game StarCraft II

<https://deepmind.com/blog/article/alphastar-mastering-real-time-strategy-game-starcraft-ii>

DeepMind and Blizzard open StarCraft II as an AI research environment

<https://deepmind.com/blog/announcements/deepmind-and-blizzard-open-starcraft-ii-ai-research-environment>

PySC2 - StarCraft II Learning Environment

<https://github.com/deepmind/pysc2>

AlphaStar: An Evolutionary Computation Perspective

<https://arxiv.org/abs/1902.01724>

StarCraft II: A New Challenge for Reinforcement Learning

<https://arxiv.org/abs/1708.04782>

3.11.2 Dota2

OpenAI five

<https://openai.com/blog/openai-five/>

Automatic player identification in Dota 2

<https://arxiv.org/abs/2008.12401>

Long-term planning and situational awareness in OpenAI five

<https://arxiv.org/abs/1912.06721>

Neural network surgery with sets

<https://arxiv.org/abs/1912.06719>

Dota 2 with large scale deep reinforcement learning

<https://arxiv.org/abs/1912.06680>

Time to die: death prediction in Dota 2 using deep learning

<https://arxiv.org/abs/1906.03939>

4 What To Turn In Via eLearning

- Upload individual files, not a zip file
- Don't write any comments in eLearning
- A pdf file of your slides
- Any optional material you'd like to include (e.g., if you do some work in PyTorch related to the project and want to share the code / results); the more impressed I am by the optional material the more additional points you can earn